

NLP Lab 8

Dixitha Kasturi
dkasturi@syr.edu

Topic: Continuation of classification and feature exploration

I explored both Option 1 and Option 2.

Option 1:

Accuracy with stop words: 0.76

Accuracy with stop words removal: 0.788

The accuracy did increase, which in a way showed that the stop words were indeed contributing negatively to the model's overall performance. I removed the stop words in the NOT_featuresets function

Top 30 features with stop words:

Most Informative Features

V_engrossing = True	pos : neg	=	19.5 : 1.0
V_stupid = True	neg : pos	=	17.8 : 1.0
V_mediocre = True	neg : pos	=	17.1 : 1.0
V_generic = True	neg : pos	=	15.1 : 1.0
V_inventive = True	pos : neg	=	14.9 : 1.0
V_routine = True	neg : pos	=	14.4 : 1.0
V_boring = True	neg : pos	=	14.2 : 1.0
V_flat = True	neg : pos	=	13.5 : 1.0
V_unique = True	pos : neg	=	12.9 : 1.0
V_refreshingly = True	pos : neg	=	12.9 : 1.0
V_refreshing = True	pos : neg	=	12.9 : 1.0
V_wonderful = True	pos : neg	=	12.5 : 1.0
V_90 = True	neg : pos	=	12.4 : 1.0
V_warm = True	pos : neg	=	11.3 : 1.0
V_stale = True	neg : pos	=	11.1 : 1.0
V_mesmerizing = True	pos : neg	=	10.9 : 1.0
V_car = True	neg : pos	=	10.4 : 1.0
V_mindless = True	neg : pos	=	10.4 : 1.0
V_dull = True	neg : pos	=	10.3 : 1.0
V_quietly = True	pos : neg	=	10.3 : 1.0
V_captures = True	pos : neg	=	10.1 : 1.0
V_powerful = True	pos : neg	=	10.0 : 1.0
V_annoying = True	neg : pos	=	9.7 : 1.0
V_provides = True	pos : neg	=	9.7 : 1.0
V_chilling = True	pos : neg	=	9.6 : 1.0
V_waste = True	neg : pos	=	9.5 : 1.0
V_supposed = True	neg : pos	=	9.1 : 1.0
V_tiresome = True	neg : pos	=	9.1 : 1.0
V_meandering = True	neg : pos	=	9.1 : 1.0
V_unexpected = True	pos : neg	=	8.9 : 1.0

Top 30 features after stop words removal:

Most Informative Features

V_engrossing = True	pos : neg =	19.7 : 1.0
V_routine = True	neg : pos =	14.9 : 1.0
V_inventive = True	pos : neg =	14.4 : 1.0
V_generic = True	neg : pos =	13.6 : 1.0
V_90 = True	neg : pos =	13.6 : 1.0
V_mediocre = True	neg : pos =	13.6 : 1.0
V_flat = True	neg : pos =	13.4 : 1.0
V_absorbing = True	pos : neg =	13.0 : 1.0
V_refreshing = True	pos : neg =	13.0 : 1.0
V_intimate = True	pos : neg =	13.0 : 1.0
V_NOTenough = True	neg : pos =	12.3 : 1.0
V_warm = True	pos : neg =	12.2 : 1.0
V_wonderful = True	pos : neg =	12.2 : 1.0
V_boring = True	neg : pos =	12.1 : 1.0
V_dull = True	neg : pos =	12.0 : 1.0
V_refreshingly = True	pos : neg =	11.7 : 1.0
V_stupid = True	neg : pos =	11.0 : 1.0
V_tv = True	neg : pos =	10.7 : 1.0
V_beauty = True	pos : neg =	10.6 : 1.0
V_provides = True	pos : neg =	10.6 : 1.0
V_thin = True	neg : pos =	10.6 : 1.0
V_touching = True	pos : neg =	10.5 : 1.0
V_extraordinary = True	pos : neg =	10.4 : 1.0
V_stale = True	neg : pos =	10.3 : 1.0
V_powerful = True	pos : neg =	9.9 : 1.0
V_captures = True	pos : neg =	9.8 : 1.0
V_tiresome = True	neg : pos =	9.6 : 1.0
V_unless = True	neg : pos =	9.6 : 1.0
V_flawed = True	pos : neg =	9.4 : 1.0
V_document = True	pos : neg =	9.0 : 1.0

Option 2:

When the given function was used with 4 numeric features the following accuracy was obtained

```
# retrain the classifier using these features
train_set, test_set = SL_featuresets[1000:], SL_featuresets[:1000]
classifier = nltk.NaiveBayesClassifier.train(train_set)
nltk.classify.accuracy(classifier, test_set)
```

0.746

When the given function was used with 1 numeric feature the following accuracy was obtained

```
def SL_features2(document, word_features, SL):
    document_words = set(document)
    features = {}
    for word in word_features:
        features['V_{}'.format(word)] = (word in document_words)
    # count variables for negative and positive
    weakPos = 0
    strongPos = 0
    weakNeg = 0
    strongNeg = 0
    for word in document_words:
        if word in SL:
            strength, posTag, isStemmed, polarity = SL[word]
            if strength == 'weaksubj' and polarity == 'positive':
                weakPos += 1
            if strength == 'strongsubj' and polarity == 'positive':
                strongPos += 1
            if strength == 'weaksubj' and polarity == 'negative':
                weakNeg += 1
            if strength == 'strongsubj' and polarity == 'negative':
                strongNeg += 1
            negcount = weakPos + strongPos
            poscount = weakNeg + strongNeg
            features['totcount'] = negcount - poscount
    return features
```

```
train_set2, test_set2 = SL_featuresets2[1000:], SL_featuresets2[:1000]
classifier2 = nltk.NaiveBayesClassifier.train(train_set2)
nltk.classify.accuracy(classifier2, test_set2)
```

0.758

```
def SL_features3(document, word_features, SL):
    document_words = set(document)
    features = {}
    for word in word_features:
        features['V_{}'.format(word)] = (word in document_words)
    # count variables for negative and positive
    negcount = 0
    poscount = 0
    for word in document_words:
        if word in SL:
            strength, posTag, isStemmed, polarity = SL[word]
            if polarity == 'positive':
                poscount += 1
            if polarity == 'negative':
                negcount += 1
            features['totcount'] = negcount - poscount
    return features
```

```
train_set3, test_set3 = SL_featuresets3[1000:], SL_featuresets3[:1000]
classifier3 = nltk.NaiveBayesClassifier.train(train_set3)
nltk.classify.accuracy(classifier3, test_set3)
```

0.758

The accuracy increased when only 1 numeric feature was considered. I wrote 2 versions of the same action to verify if my approach was correct.